

Last Name:Jengirapas

First Name:Natcha

Student ID:85939811

1.

a) CQL Query:

```
DESC "Hoofers";
```

b) Result:

```
CREATE KEYSPACE "Hoofers" WITH replication = {'class': 'NetworkTopologyStrategy', 'us-west-2': '3'} AND durable_writes = true;
CREATE TABLE "Hoofers".boats (
  bid int PRIMARY KEY,
  bname text,
  color text
) WITH additional_write_policy = '99PERCENTILE'
  AND bloom_filter_fp_chance = 0.01
  AND caching = {'keys': 'ALL', 'rows_per_partition': 'NONE'}
  AND comment = ''
  AND compaction = {'class': 'org.apache.cassandra.db.compaction.UnifiedCompactionStrategy', 'log_all': 'true', 'num_shards': '128'}
  AND compression = {'chunk_length_in_kb': '64', 'class': 'org.apache.cassandra.io.compress.LZ4Compressor'}
  AND crc_check_chance = 1.0
  AND default_time_to_live = 0
  AND gc_grace_seconds = 864000
  AND max_index_interval = 2048
  AND memtable_flush_period_in_ms = 0
  AND min_index_interval = 128
  AND read_repair = 'BLOCKING'
  AND speculative_retry = '99PERCENTILE';
```

c) Answers:

- The keyspace maintains 3 copies of the data
- The cloud region is Us-west-2
- Greater than 3

2.

a) CQL CREATE Statements:

```
CREATE TABLE Customers (  
    customer_id VARCHAR,  
    PRIMARY KEY(customer_id)  
);  
  
CREATE TABLE Orders (  
    order_id VARCHAR,  
    total_price DECIMAL,  
    time_placed TIMESTAMP,  
    pickup_time TIMESTAMP,  
    customer_id VARCHAR,  
    shopper_id VARCHAR,  
    state VARCHAR,  
    license_plate VARCHAR,  
    store_id VARCHAR,  
    time_fulfilled TIMESTAMP,  
    PRIMARY KEY (order_id)  
);  
  
CREATE TABLE OrderItems (  
    item_id VARCHAR,  
    order_id VARCHAR,  
    product_id VARCHAR,  
    qty INT,  
    selling_price DECIMAL,  
    PRIMARY KEY(item_id, order_id)  
);  
  
CREATE TABLE Products (  
    product_id VARCHAR,  
    category VARCHAR,  
    name VARCHAR,  
    description VARCHAR,  
    list_price DECIMAL,  
    PRIMARY KEY (product_id)  
);
```

3.

a) PostgreSQL COPY commands:

```
copy Customers to '/Users/rew/Desktop/School/CS 122D/HW/HW2/CSV
files/customers.csv' DELIMITER ',' CSV HEADER;
copy Products to '/Users/rew/Desktop/School/CS 122D/HW/HW2/CSV
files/products.csv' DELIMITER ',' CSV HEADER;
copy Orders to '/Users/rew/Desktop/School/CS 122D/HW/HW2/CSV
files/orders.csv' DELIMITER ',' CSV HEADER;
copy Orderitems to '/Users/rew/Desktop/School/CS 122D/HW/HW2/CSV
files/orderitems.csv' DELIMITER ',' CSV HEADER;
```

4.

a) First CQL Query:

```
SELECT name, list_price FROM Products WHERE category = 'Meat & Seafood'
LIMIT 10;
```

b) Result:

```
token@cqlsh:ShopALot> SELECT name, list_price FROM Products WHERE category = 'Meat & Seafood' LIMIT 10;
InvalidRequest: Error from server: code=2200 [Invalid query] message="Cannot execute this query as it might involve data filtering and thus may have unpredictable performance. If you want to execute this query despite the performance unpredictability, use ALLOW FILTERING"
token@cqlsh:ShopALot>
```

c) Modified CQL Query:

```
SELECT name, list_price FROM Products WHERE category = 'Meat & Seafood'
LIMIT 10 ALLOW FILTERING;
```

d) Result:

| name | list_price |
|---|------------|
| USDA Choice Beef Top Loin New York Strip Steak Bone In Value Pack - 3.5 Lbs. (approx. weight) | 34.97 |
| waterfront BISTRO Salmon Fillets Wild Alaskan Pink Boneless & Skin On - 32 Oz | null |
| Signature Farms Boneless Skinless Chicken Thighs Value Pack - 3 Lbs. | 13.47 |
| Aidells Smoked Chicken Sausage Links Chicken & Apple 4 Count - 12 Oz | 6.49 |
| Foster Farms Fresh & Natural Whole Chicken - 3.5 Lbs. | 6.97 |
| USDA Choice Beef Filet Mignon Steak Tenderloin - 1 Lb. | 19.99 |
| Jennie-O Ground Turkey 93% Lean 7% Fat - 16 Oz. | 5.99 |
| Jimmy Dean Fully Cooked Original Pork Sausage Links 12 Count - 9.6 Oz | 5.99 |
| Signature Farms Boneless Skinless Chicken Breasts Value Pack - 3 Lbs. | 13.47 |
| Tyson Grilled & Ready Fully Cooked Grilled Chicken Breast Strips - 22 Oz | 8.99 |

(10 rows)
token@cqlsh:ShopALot>

5.

a) CQL Create Statement:

```
CREATE TABLE Productsq5 (  
    product_id VARCHAR,  
    category VARCHAR,  
    name VARCHAR,  
    description VARCHAR,  
    list_price DECIMAL,  
    PRIMARY KEY (category, product_id)  
);
```

b) CQL Query:

```
SELECT name, list_price  
FROM Productsq5  
WHERE category = 'Meat & Seafood'  
LIMIT 10;
```

c) Result:

```
token@cqlsh:ShopALot> SELECT name, list_price  
... FROM Productsq5  
... WHERE category = 'Meat & Seafood'  
... LIMIT 10;
```

| name | list_price |
|--|------------|
| Ground Beef 80% Lean 20% Fat - 1.25 Lbs. | 4.99 |
| USDA Choice Beef Filet Mignon Steak Tenderloin - 1 Lb. | 19.99 |
| Jimmy Dean Fully Cooked Turkey Sausage Links 12 Count - 9.6 Oz | 5.99 |
| Seafood Counter Fish Salmon Fresh Atlantic Salmon Fillet Color Added - 1.00 LB | 9.99 |
| Tyson Fully Cooked Breaded Chicken Nuggets - 32 Oz | 7.99 |
| Tyson Grilled & Ready Fully Cooked Grilled Chicken Breast Strips - 22 Oz | 8.99 |
| Tyson Anytizers Buffalo Boneless Chicken Bites - 24 Oz | 0.99 |
| Signature Farms Boneless Skinless Chicken Breasts Value Pack - 3 Lbs. | 13.47 |
| Jimmy Dean Fully Cooked Original Pork Sausage Links 12 Count - 9.6 Oz | 5.99 |
| Signature Farms Frozen Boneless Skinless Chicken Breasts - 40 Oz. | 8.99 |

(10 rows)

d) Explanation:

Partition keys allow Cassandra to use hashing to find where that key-values would be. Having the category as a partitioning key, Cassandra knows which node it has to retrieve data from, which it is allowed.

If I only included the category as a primary key, the table will only have 16 rows representing every 16 categories since the partition key is required to be unique. However, if I included the category as partition key and product_id as clustering key, the table will contain all the data such that each category will have all the sorted product_id in ascending order of that category.

6.

a) CQL Query:

```
SELECT name, list_price
FROM Productsq6
WHERE category = 'Meat & Seafood'
ORDER BY list_price desc
LIMIT 10;
```

b) CQL CREATE Statement:

```
CREATE TABLE Productsq6 (
    product_id VARCHAR,
    category VARCHAR,
    name VARCHAR,
    description VARCHAR,
    list_price DECIMAL,
    PRIMARY KEY (category, list_price, product_id)
);
```

c) Results:

```
(10 rows)
token@cqlsh:ShopALot> SELECT name, list_price
... FROM Productsq6
... WHERE category = 'Meat & Seafood'
... ORDER BY list_price desc
... LIMIT 10;
```

| name | list_price |
|---|------------|
| USDA Choice Beef Ribeye Roast Bone In - 6 Lbs. (approx. weight) | 59.94 |
| USDA Choice Beef Top Loin New York Strip Steak Bone In Value Pack - 3.5 Lbs. (approx. weight) | 34.97 |
| Signature Farms Beef Corned Beef Brisket Flat Cut - 3.50 LB | 22.72 |
| USDA Choice Beef Filet Mignon Steak Tenderloin - 1 Lb. | 19.99 |
| USDA Choice Beef Boneless Chuck Roast - 3 Lbs. | 17.97 |
| Signature Farms Boneless Skinless Chicken Thighs Value Pack - 3 Lbs. | 13.47 |
| Signature Farms Boneless Skinless Chicken Breasts Value Pack - 3 Lbs. | 13.47 |
| Open Nature Frozen Boneless Skinless Chicken Breasts - 36 Oz. | 11.99 |
| Seafood Counter Fish Salmon Fresh Atlantic Salmon Fillet Color Added - 1.00 LB | 9.99 |
| Hormel Black Label Bacon Original - 16 Oz | 8.99 |

```
(10 rows)
```

d) Explanation:

Cassandra achieves performance through the use of the clustering keys to order data. Thus, returning ordered rows in a single read. The primary key is similar to the previous question. However, we included list_price as another clustering key before product_id. Hence, Cassandra would order the data by list_price first then product_id.

7.

a) CQL Create Statement:

```
CREATE TABLE Ordersq7a (  
    customer_id VARCHAR,  
    order_id VARCHAR,  
    total_price DECIMAL,  
    PRIMARY KEY(customer_id, total_price, order_id)  
) WITH CLUSTERING ORDER BY (total_price DESC, order_id ASC);
```

b) CQL Create Statement:

```
CREATE TABLE Shoppersq7b(  
    shopper_id VARCHAR,  
    order_id VARCHAR,  
    item_id VARCHAR,  
    PRIMARY KEY (shopper_id, order_id, item_id));
```

c) CQL Create Statement:

```
CREATE TABLE Ordersq7C(  
    order_id VARCHAR,  
    item_id VARCHAR,  
    name VARCHAR,  
    category VARCHAR,  
    list_price DECIMAL,  
    PRIMARY KEY (order_id, list_price, item_id));
```

d) CQL Create Statement:

```
CREATE TABLE Customersq7d(  
    customer_id VARCHAR,  
    order_id VARCHAR,  
    total_price DECIMAL,  
    time_placed TIMESTAMP,  
    PRIMARY KEY(customer_id, time_placed, order_id));
```

8.

a)

- CQL Query:

```
SELECT order_id FROM Ordersq7a WHERE Customer_id = '24590';
```

- Result:

```
order_id
-----
87J33
UDK7R
2CFUS
WPKIJ
WK0H4
A4V99
NA0TI
(7 rows)
```

b)

- CQL Query:

```
SELECT order_id, count(*) FROM Shoppersq7b WHERE shopper_id = '0JKLY'
group by order_id limit 10;
```

- Result:

```
token@cqlsh:ShopALot> SELECT order_id, count(*) FROM Shoppersq7b WHERE shopper_id = '0JKLY' group by order_id limit 10;
```

| order_id | count |
|----------|-------|
| 07DY0 | 1 |
| 9MLF9 | 2 |
| A3BRA | 1 |
| NHOKA | 1 |
| SJ097 | 1 |

(5 rows)

c)

- CQL Query:

```
SELECT name, category FROM Ordersq7c WHERE order_id = '005SN';
```

- Result:

```
token@cqlsh:ShopALot> SELECT name, category FROM Ordersq7c WHERE order_id = '005SN';
```

| name | category |
|--|------------------------|
| Horizon Organic Milk Reduced Fat 2% - Half Gallon | Dairy & Eggs & Cheese |
| Larabar Food Bar Apple Pie - 5-1.6 Oz | Personal Care & Health |
| Mini Babybel Original Snack Cheese - 10 Count - 7.5 oz | Deli |
| Egglands Best Eggs Large Grade A - 18 Count | Dairy & Eggs & Cheese |

(4 rows)

d)

- CQL Query:

```
SELECT SUM(total_price) FROM Customersq7d WHERE customer_id = '32976' and  
time_placed >= '2020-03-01 00:00:00' and time_placed <= '2020-09-01  
00:00:00';
```

- Result:

```
system.sum(total_price)  
-----  
614.47  
  
(1 rows)
```


9.

a) CQL INSERT statements:

```
-- Tables from Part 2
INSERT INTO orders (order_id, total_price, pickup_time, customer_id,
shopper_id, state, license_plate, store_id, time_fulfilled) VALUES
('12MDAE', 7.14, '2021-04-10T21:01:45.000Z', '24590', 'MQD30', 'CA', 'AKM
554', 'A7ZNF', '2021-04-10T23:20:56.000Z');

INSERT INTO orderItems (item_id, order_id, product_id, qty, selling_price)
VALUES ('9B317', '12MDAE', 'GMG05', 3, 2.38);

-- Tables from Part 7
INSERT INTO ordersq7a (customer_id, order_id, total_price) VALUES ('24590',
'12MDAE', 7.14);

INSERT INTO shoppersq7b (shopper_id, order_id, item_id) VALUES ('MQD30',
'12MDAE', '9B317');

INSERT INTO ordersq7c (order_id, item_id, name, category, list_price)
VALUES ( '12MDAE', '9B317', 'Green Giant Corn Whole Kernel Sweet - 15.25
Oz', 'Canned Goods & Soups', 1.59);

INSERT INTO customersq7d (customer_id, order_id, total_price, time_placed)
VALUES ('24590', '9B317', 7.14, '2021-04-10T19:30:24.000Z');
```

10. [Extra Credit]

Python script:

```
from cassandra.cluster import Cluster
from cassandra.auth import PlainTextAuthProvider
from datetime import datetime

# Connecting to the database: code from DataStax Astra
def connect_cassandra():
    cloud_config = {
        'secure_connect_bundle': "/Users/rew/Downloads/secure-connect-cs122d-spring.zip"
    }
    auth_provider = PlainTextAuthProvider('rsvrHPKpQDmYhjDjhBIISHJZ',
    'ENo2M,rKX41d1PzQv2h1SvhmZex9I1Ft7d+DISg952d18TfRhlinySSPuKMvm3ZLMzaXvULKp00uWGn.LBAwn.BxuU
qh8w0d86e0C8rEUPYQzHAoG.mg8E0WvL97T2h')
    cluster = Cluster(cloud=cloud_config, auth_provider=auth_provider)
    global session
    session = cluster.connect()

    row = session.execute("select release_version from system.local").one()

    if row:
        print("SUCCESSFULLY CONNECTED!")
    else:
        print("An error occurred.")

# Getting all table names and column names in a keyspace
def get_tables(keyspace):
    tables = dict()
    tempTables = session.execute(f"SELECT * FROM system_schema.tables WHERE keyspace_name = '{keyspace}';")

    for table in tempTables:
        tempColumns = session.execute(f"SELECT * FROM system_schema.columns WHERE keyspace_name = '{keyspace}' AND table_name = '{table.table_name}';")
        column = {i.column_name: 0 for i in tempColumns}
        tables[table.table_name] = column

    return tables

# Creating an insert command based on the argument dictionary
def create_insert(tables, table_name, argDict):
    text = ""

    for column in tables[table_name]:
        if column in argDict:
            text += f"{column}\": \"{argDict[column]}\", "
        elif column in ["name", "category", "list_price"]:
            product_id = argDict["orderItems"][0]["product_id"]
```

```

        lookup = session.execute(f"SELECT {column} FROM \"ShopALot\".Products WHERE
product_id = '{product_id}'")
        text += f"\"{column}\": \"{lookup.one()[0]}\", "
    elif column in ["item_id", "qty", "selling_price", "product_id"]:
        orderItem = argDict["orderItems"][0][column]
        text += f"\"{column}\": \"{orderItem}\", "

text = f"INSERT INTO \"ShopALot\".{table_name} json '{{{text[:-2]}}}'";

insert = session.prepare(text)
results = session.execute(insert, argDict)

def insert_new_order(tables, infoDict):
    for table in tables:
        # Check if the input table needs to add new order
        if "order_id" in tables[table]:
            create_insert(tables, table, infoDict)

def main():
    keyspace = "ShopALot"
    connect_cassandra()
    new_order = {"order_id": "WEQ174", "total_price": 36.47, "time_placed":
"2021-03-29T15:03:20.000Z",
                "pickup_time": "2021-03-23T17:54:21.000Z", "customer_id": "6Z53Z",
"shopper_id": "MQD30",
                "state": "WV", "license_plate": "0031", "store_id": "ZU9IP",
"time_fulfilled": "2021-03-23T22:43:12.000Z",
                "orderItems": [{"item_id": "P012C", "qty": "7", "selling_price": 5.21,
"product_id": "GMG05"}]}

    tables = get_tables(keyspace)
    insert_new_order(tables, new_order)

if __name__ == "__main__":
    main();

```

*Query to show that the new order has been added to the table.

```

token@cqlsh:ShopALot> select * from customersq7d where customer_id = '6Z53Z';

```

| customer_id | time_placed | order_id | total_price |
|-------------|---------------------------------|----------|-------------|
| 6Z53Z | 2019-04-10 22:04:42.000000+0000 | QMQB5 | 47.78 |
| 6Z53Z | 2020-01-14 04:34:01.000000+0000 | ZRMH6 | 242.47 |
| 6Z53Z | 2020-01-26 20:49:59.000000+0000 | WXW4X | 30.36 |
| 6Z53Z | 2020-06-04 22:46:45.000000+0000 | DP8TJ | 27.00 |
| 6Z53Z | 2020-06-15 00:38:03.000000+0000 | NEAKE | 7.90 |
| 6Z53Z | 2020-07-20 23:36:20.000000+0000 | MU8PA | 7.90 |
| 6Z53Z | 2020-07-26 11:22:25.000000+0000 | BCPP2 | 15.00 |
| 6Z53Z | 2021-03-29 15:03:20.000000+0000 | WEQ174 | 36.47 |

```

(8 rows)
token@cqlsh:ShopALot>

```